



RISC-V Matrix Specification

Xin Ouyang, Zhiqiang Liu

Version 0.1, 9/2022: This document is in development.

Table of Contents

Preamble	1
Copyright and license information	2
Contributors	3
1. Introduction	4
2. Implementation-defined Constant Parameters	5
3. Programmer's Model	6
3.1. Matrix Registers	6
3.2. Matrix type register, mtype	7
3.3. Matrix tile configure registers, tile_m/tile_k/tile_n	7
4. Instructions	8
4.1. Instruction Formats	8
4.2. Configuration-Setting Instructions	8
4.2.1. mtype encoding	9
4.2.2. ATM/ATK/ATN encoding	10
4.2.3. Constraints on Setting tile_m/tile_k/tile_n	11
4.3. Load and Store Instructions	12
4.3.1. Load instructions	12
4.3.2. Store instructions	13
4.4. Arithmetic Instructions	14
4.4.1. Matrix Multiplication Instructions	14
4.4.2. Element-Wise Add/Sub/Multiply Instructions	15
4.4.3. Type-Convert Instructions	16
4.5. Instruction Listing	16
5. Intrinsic Examples	21
5.1. Matrix multiplication	21
5.2. Matrix multiplication with left matrix transposed	21
5.3. Matrix transpose without multiplication	22
6. Standard Matrix Extensions	23
6.1. Zmv: Matrix for Vector operations	23
6.1.1. Load Instructions	23
6.1.2. Store Instructions	24
6.1.3. Data Move Instructions	24
6.1.4. Matrix element-wise multiply	26
6.1.5. Instruction Listing	26
6.1.6. Intrinsic Examples: Matrix multiplication fused with element-wise vector operation ..	29
6.2. Zmbf16: Matrix Bfloat16(BF16) Extension	29

6.3. Zmtf32: Matrix TensorFloat-32(TF32) Extension	30
6.4. Zmic: Im2col Matrix Multiplication Extension	31
Bibliography	?

Preamble



This document is in the [Development state](#)

Assume everything can change. This draft specification will change before being accepted as standard, so implementations made to this draft specification will likely not conform to the future standard.

Copyright and license information

This specification is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at creativecommons.org/licenses/by/4.0/.

Copyright 2022 by Streamcomputing Corp.

Contributors

This RISC-V specification has been contributed to directly or indirectly by:

- Kening Zhang
- Kun Hu
- Hui Yao
- **Your name here**

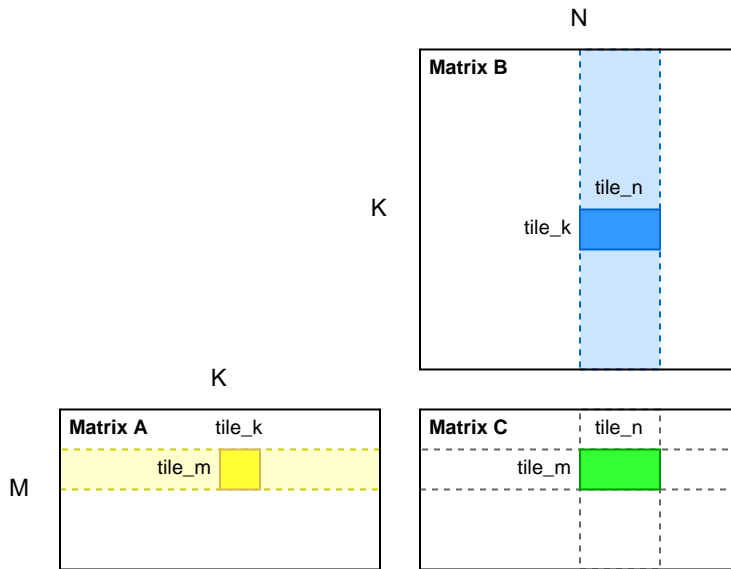
We will be very grateful to the huge number of other people who will have helped to improve this specification through their comments, reviews, feedback and questions.

Chapter 1. Introduction

This document describes the matrix extension for RISC-V.

Matrix extension implement matrix multiplications by partitioning the input and output matrix into tiles, which are then stored to matrix registers.

Tile size usually refers to the dimensions of these tiles. For the operation $C = AB$ in figure below, the tile size of C is $\text{tile_m} \times \text{tile_n}$, the tile size of A is $\text{tile_m} \times \text{tile_k}$ and the tile size of B is $\text{tile_k} \times \text{tile_n}$.



Each matrix multiplication instruction computes its output tile by stepping through the K dimension in tiles, loading the required values from the A and B matrices, and multiplying and accumulating them into the output.

Matrix extension is strongly inspired by the RISC-V Vector "V" extension.

Chapter 2. Implementation-defined Constant Parameters

Each hart supporting a matrix extension defines three parameters:

1. The maximum size in bits of a matrix element that any operation can produce or consume, $ELEN \geq 8$, which must be a power of 2.
2. The number of bits in a single matrix tile register, $MLEN \geq ELEN$, which must be a power of 2, and must be no greater than 2^{32} .
3. The number of bits in a row of a single matrix tile register, $RLEN \geq ELEN$, $RLEN \leq MLEN$, which must be a power of 2, and must be no greater than 2^{16} .

Chapter 3. Programmer's Model

The matrix extension adds 10 matrix tile registers, and 5 unprivileged CSRs (tile_m, tile_n, tile_k, mtype, mlenb) to the base scalar RISC-V ISA.

Table 1. New matrix CSRs

Address	Privilege	Name	Description
0xC40	URO	tile_m	Tile length in m direction
0xC41	URO	tile_n	Tile length in n direction
0xC42	URO	tile_k	Tile length in k direction
0xC43	URO	mtype	Matrix tile data type register
0xC44	URO	mlenb	MLEN/8 (matrix tile register length in bytes)

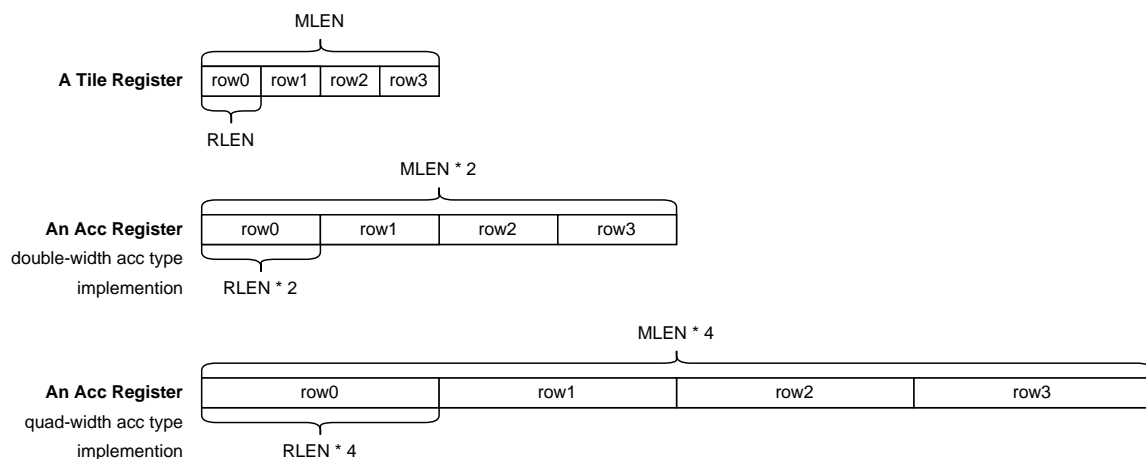
3.1. Matrix Registers

The matrix extension adds 10 architectural matrix registers, 8 **Tile Registers**(tr0-tr7) for tiles of input matrices, and 2 **Accumulator Registers**(acc0-acc1) for tiles of output matrices.

A **Tile Register** has a fixed MLEN bits of state.

If a implementation supports using a accumulator data type which is quad-width of the input data type, each **Accumulator Register** would have a MLEN*4 bits of state.

Otherwise, each **Accumulator Register** has a MLEN*2 bits of state, to support a double-width accumulator data type.



3.2. Matrix type register, `mtype`

The read-only XLEN-wide *matrix type* CSR, `mtype` provides the default type used to interpret the contents of the matrix register file, and can only be updated by `msettype{i}` instructions. The matrix type determines the organization of elements in each matrix register.



Allowing updates only via the `msettype{i}` instructions simplifies maintenance of the `mtype` register state.

The `mtype` register has five fields, `mill`, `maccq`, and `msew[2:0]`. Bits `mtype[XLEN-2:4]` should be written with zero, and non-zero values in this field are reserved.

Table 2. `mtype` register layout

Bits	Name	Description
XLEN-1	<code>mill</code>	Illegal value if set
XLEN-2:4	0	Reserved if non-zero
3	<code>maccq</code>	Support quad-width accumulator element
2:0	<code>msew[2:0]</code>	Selected element width (SEW) setting

3.3. Matrix tile configure registers, `tile_m/tile_k/tile_n`

The XLEN-bit-wide read-only `tile_m/tile_k/tile_n` CSRs can only be updated by the `msettile[m|k|n]{i}` instructions. The registers hold 3 unsigned integers specifying the tile shapes for tiled matrix.

Chapter 4. Instructions

4.1. Instruction Formats

The instructions in the matrix extension use a new major opcode (1110111, which inst[6:5]=11, inst[4:2]=101 is reserved in RISC-V opcode map).

This instruction formats are list below.

Configuration instructions, funct3 = 111

31 28	27 20	19 15	14 12	11 7	6 0
funct4	imm13		funct3	rd	1110111
funct4	00000000	rs1	funct3	rd	1110111

Load & Store instructions, eew = 000 - 011

31 26	25	24 20	19 15	14 12	11 7	6 0
funct6	ls	rs2	rs1	eew	md	1110111

Arithmetic & Type-Convert instructions, funct3 = 110

31 26	25	24 20	19 15	14 12	11 7	6 0
funct6	fp	ms2	ms1	funct3	md	1110111

Data Move instructions, funct3 = 101

31 26	25	24 20	19 15	14 12	11 7	6 0
funct6	di	rs2	ms1	funct3	md	1110111

4.2. Configuration-Setting Instructions

Due to hardware resource constraints, one of the common ways to handle large-sized matrix multiplications is "tiling", where each iteration of the loop processes a subset of elements, and then continues to iterate until all elements are processed. The Matrix extension provides direct, portable support for this approach.

The block processing of matrix multiplication requires three levels of loops to iterate in the direction of the number of rows of the left matrix (m), the number of columns of the left matrix(k, also the number of rows of the right matrix), and the number of columns of the right matrix(n), given by the application.

The shapes of the matrix tiles to be processed, m (application tile length m or **ATM**), k (**ATK**), n (**ATN**), is used as candidates for **tile_m/tile_k/tile_n**. Based on microarchitecture implementation and **mtype** setting, hardware returns a new **tile_m/tile_k/tile_n** value via a general purpose register (usually smaller), also stored in **tile_m/tile_k/tile_n** CSR, which is the shape of tile per iteration handled by hardware.

For a simple matrix multiplication example, check out the Section Intrinsic Example, which describes how the code keeps track of the matrices processed by the hardware each iteration.

A set of instructions is provided to allow rapid configuration of the values in **tile_*** and **mtype** to match application needs.

The **msettype{i}** instructions set the **mtype** CSR based on their arguments, and write the new value of **mtype** into **rd**.

msettypei rd, mtypei	# rd = new mtype, mtypei = new mtype setting
msettype rd, rs1	# rd = new mtype, rs1 = new mtype value

The **msettile[m|k|n]{i}** instructions set the **tile_m/tile_k/tile_n** CSRs based on their arguments, and write the new value into **rd**.

msettilemi rd, mleni	# rd = new tile_m, mleni = ATM
msettilem rd, rs1	# rd = new tile_m, rs1 = ATM
msettileki rd, mleni	# rd = new tile_k, mleni = ATN
msettilek rd, rs1	# rd = new tile_k, rs1 = ATN
msettileni rd, mleni	# rd = new tile_n, mleni = ATK
msettilen rd, rs1	# rd = new tile_n, rs1 = ATK

4.2.1. **mtype** encoding

The **mtype** register has five fields, **mill**, **maccq**, and **msew[2:0]**. Bits **mtype[XLEN-2:4]** should be written with zero, and non-zero values in this field are reserved.

Table 3. **mtype** register layout

Bits	Name	Description
XLEN-1	mill	Illegal value if set
XLEN-2:4	0	Reserved if non-zero
3	maccq	Support quad-width accumulator element
2:0	msew[2:0]	Selected element width (SEW) setting

The new **mtype** value is encoded in the immediate fields of **msettypei**, and in the **rs1** register for

msettype.

Suggested assembler names used for msettypei mtypei immediate

```
e8    # SEW = 8b
e16   # SEW = 16b
e32   # SEW = 32b
e64   # SEW = 64b
```

```
accq  # support 32-bit accumulator element
```

Examples:

```
msettypei t0, e8           # SEW = 8
msettypei t0, e32          # SEW = 16

msettypei t0, e8, accq     # SEW = 8, support 32-bit accumulator element
```

4.2.2. ATM/ATK/ATN encoding

There are three values, **TMMAX**, **TKMAX**, **TNMAX**, represents the maximum shapes of the matrix tiles could be stored in matrix registers, that can be operated on with a single matrix instruction given the current SEW settings as shown below.

- $TMMAX = MLEN / RLEN$
- $TKMAX = \min(MLEN / RLEN, RLEN / SEW)$
- $TNMAX = RLEN / SEW$

For examples, if $MLEN=256$, $RLEN=64$, **TMMAX**, **TKMAX**, **TNMAX** values are shown below.

```
SEW=8,  TMMAX=4, TKMAX=4, TNMAX=8      # 4x4x8 8bit matmul
SEW=16, TMMAX=4, TKMAX=4, TNMAX=4,     # 4x4x4 16bit matmul
SEW=32, TMMAX=4, TKMAX=2, TNMAX=2,     # 4x2x2 32bit matmul
```

The new tile shape settings are based on **ATM/ATK/ATN** values, which for **msettile[m|k|n]{i}** is encoded in the rs1 and rd fields.

rd	rs1	ATM/ATK/ATN value	Effect on tile_m/tile_k/tile_n
-	!x0	Value in x[rs1]	Normal tiling
!x0	x0	~0	Set tile_m/tile_k/tile_n to TMMAX/TKMAX/TNMAX

x0	x0	Value in <code>tile_m/tile_k/tile_n</code>	Keep existing <code>tile_m/tile_k/tile_n</code> if less than <code>TMMAX/TKMAX/TNMAX</code>
----	----	--	---

For the `msettile[m|k|n]{i}` instructions, the `ATM/ATK/ATN` is encoded as a 13-bit zero-extended immediate in the `rs1`.

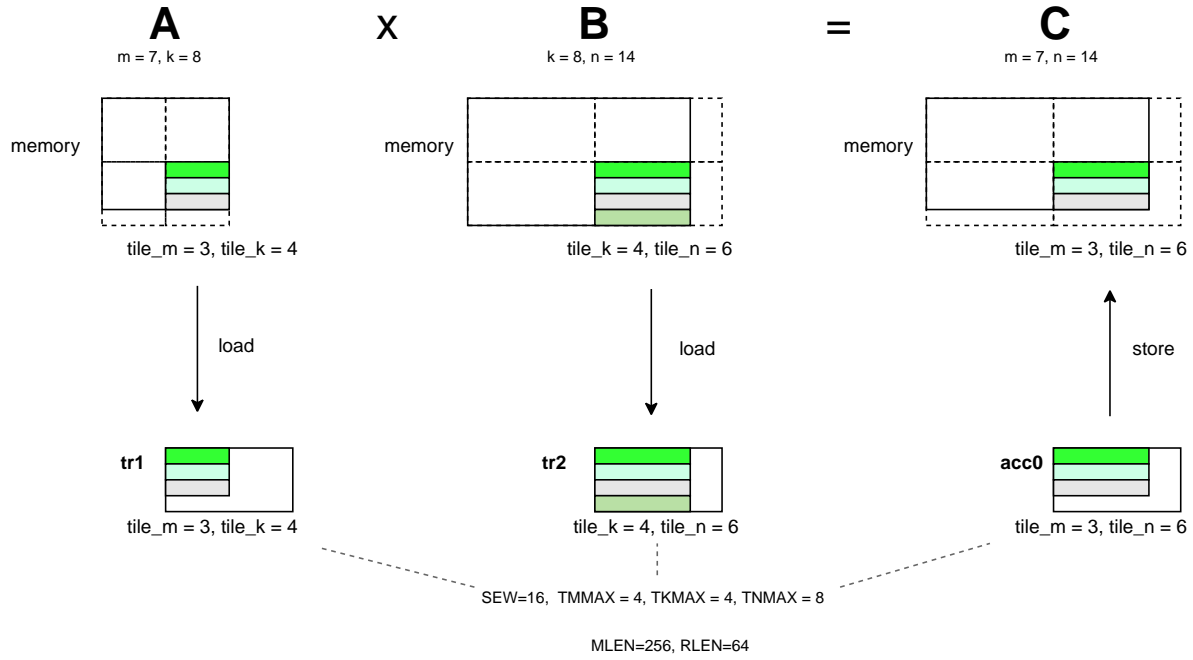
4.2.3. Constraints on Setting `tile_m/tile_k/tile_n`

The `msettile[m|k|n]{i}` instructions first set `TMMAX/TKMAX/TNMAX` according to the `mtype` CSR, then set `tile_m/tile_k/tile_n` obeying the following constraints (use `tile_m&ATM&TMMAX` as example, same to `tile_k&ATK&TKMAX` and `tile_n&ATN&TNMAX`):

1. `tile_m = ATM` if `ATM <= TMMAX`
2. `ceil(ATM / 2) <= tile_m <= TMMAX` if `ATM < (2 * TMMAX)`
3. `tile_m = TMMAX` if `ATM >= (2 * TMMAX)`
4. Deterministic on any given implementation for same input `ATM` and `TMMAX` values
5. These specific properties follow from the prior rules:
 - a. `tile_m = 0` if `ATM = 0`
 - b. `tile_m > 0` if `ATM > 0`
 - c. `tile_m <= TMMAX`
 - d. `tile_m <= ATM`
 - e. a value read from `tile_m` when used as the `ATM` argument to `msettile[m|k|n]{i}` results in the same value in `tile_m`, provided the resultant `TMMAX` equals the value of `TMMAX` at the time that `tile_m` was read.

Continue to use `MLEN=256`, `RLEN=64` as a example. When `SEW=16`, `TMMAX=4`, `TKMAX=4`, `TNMAX=8`.

If `A` is a 7 x 8 matrix and `B` is a 8 x 14 matrix, we could get `tile_m/tile_k/tile_n` values as show below, in the last loop of tiling.



4.3. Load and Store Instructions

4.3.1. Load instructions

Load a matrix tile from memory.



For **mla***, **msa***, **mlb*** or **msb*** instructions, source and destination operands are **tile registers**(**tr0-tr7**). And for **mlc***, **msc*** instructions, source and destination operands are **accumulator registers**(**acc0-acc1**)

```
# md destination, rs1 base address, rs2 row byte stride

# for left matrix, a
# tile size = tile_m * tile_k
mlae8.m md, (rs1), rs2 # 8-bit tile load
mlae16.m md, (rs1), rs2 # 16-bit tile load
mlae32.m md, (rs1), rs2 # 32-bit tile load
mlae64.m md, (rs1), rs2 # 64-bit tile load

# for right matrix, b
# tile size = tile_k * tile_n
mlbe8.m md, (rs1), rs2 # 8-bit tile load
mlbe16.m md, (rs1), rs2 # 16-bit tile load
mlbe32.m md, (rs1), rs2 # 32-bit tile load
mlbe64.m md, (rs1), rs2 # 64-bit tile load
```

```

# for output matrix, c
# tile size = tile_m * tile_n
mlce8.m  md, (rs1), rs2 # 8-bit acc load
mlce16.m md, (rs1), rs2 # 16-bit acc load
mlce32.m md, (rs1), rs2 # 32-bit acc load
mlce64.m md, (rs1), rs2 # 64-bit acc load

```

Load a matrix tile from memory, the matrix on memory is transposed.

```

# md destination, rs1 base address, rs2 row byte stride

# for left matrix, a
# tile size = tile_k * tile_m
mlate8.m  md, (rs1), rs2 # 8-bit tile load
mlate16.m md, (rs1), rs2 # 16-bit tile load
mlate32.m md, (rs1), rs2 # 32-bit tile load
mlate64.m md, (rs1), rs2 # 64-bit tile load

# for right matrix, b
# tile size = tile_n * tile_k
mlbte8.m  md, (rs1), rs2 # 8-bit tile load
mlbte16.m md, (rs1), rs2 # 16-bit tile load
mlbte32.m md, (rs1), rs2 # 32-bit tile load
mlbte64.m md, (rs1), rs2 # 64-bit tile load

# for output matrix, c
# tile size = tile_n * tile_m
mlcte8.m  md, (rs1), rs2 # 8-bit acc load
mlcte16.m md, (rs1), rs2 # 16-bit acc load
mlcte32.m md, (rs1), rs2 # 32-bit acc load
mlcte64.m md, (rs1), rs2 # 64-bit acc load

```

4.3.2. Store instructions

Store a matrix tile to memory.

```

# ms3 store data, rs1 base address, rs2 row byte stride

# for left matrix, a
# tile size = tile_m * tile_k
msae8.m  ms3, (rs1), rs2 # 8-bit tile store
msae16.m ms3, (rs1), rs2 # 16-bit tile store
msae32.m ms3, (rs1), rs2 # 32-bit tile store
msae64.m ms3, (rs1), rs2 # 64-bit tile store

# for right matrix, b

```



```

# tile size = tile_k * tile_n
msbe8.m  ms3, (rs1), rs2 # 8-bit tile store
msbe16.m ms3, (rs1), rs2 # 16-bit tile store
msbe32.m ms3, (rs1), rs2 # 32-bit tile store
msbe64.m ms3, (rs1), rs2 # 64-bit tile store

# for output matrix, c
# tile size = tile_m * tile_n
msce8.m  ms3, (rs1), rs2 # 8-bit acc store
msce16.m ms3, (rs1), rs2 # 16-bit acc store
msce32.m ms3, (rs1), rs2 # 32-bit acc store
msce64.m ms3, (rs1), rs2 # 64-bit acc store

```

Save a matrix tile to memory, the matrix on memory is transposed.

```

# ms3 store data, rs1 base address, rs2 row byte stride

# for left matrix, a
# tile size = tile_k * tile_m
msate8.m  ms3, (rs1), rs2 # 8-bit tile store
msate16.m ms3, (rs1), rs2 # 16-bit tile store
msate32.m ms3, (rs1), rs2 # 32-bit tile store
msate64.m ms3, (rs1), rs2 # 64-bit tile store

# for right matrix, b
# tile size = tile_n * tile_k
msbte8.m  ms3, (rs1), rs2 # 8-bit tile store
msbte16.m ms3, (rs1), rs2 # 16-bit tile store
msbte32.m ms3, (rs1), rs2 # 32-bit tile store
msbte64.m ms3, (rs1), rs2 # 64-bit tile store

# for output matrix, c
# tile size = tile_n * tile_m
mscte8.m  ms3, (rs1), rs2 # 8-bit acc store
mscte16.m ms3, (rs1), rs2 # 16-bit acc store
mscte32.m ms3, (rs1), rs2 # 32-bit acc store
mscte64.m ms3, (rs1), rs2 # 64-bit acc store

```

4.4. Arithmetic Instructions

4.4.1. Matrix Multiplication Instructions

Matrix Multiplication operations take two matrix tiles of operands from matrix **tile registers** specified by **ms2** and **ms1** respectively, the output matrix tiles of operands is a matrix **accumulator register** specified by **md**.

```
# int matrix multiplication and add, md = md + ms1 * ms2
mma.mm md, ms1, ms2
mwma.mm md, ms1, ms2    # output double widen
mqma.mm md, ms1, ms2    # output quadruple widen

# float matrix multiplication and add, md = md + ms1 * ms2
mfma.mm md, ms1, ms2
mfwma.mm md, ms1, ms2   # output double widen
```

4.4.2. Element-Wise Add/Sub/Multiply Instructions

Matrix element-wise add/sub/multiply instructions.

```
# int matrix element-wise add, md[i,j] = md[i,j] + ms1[i,j]
maddc.mm md, ms1
mwaddc.mm md, ms1        # output double widen
mqaddc.mm md, ms1        # output quadruple widen

# int matrix element-wise subtract, md[i,j] = md[i,j] - ms1[i,j]
msubc.mm md, ms1
mwsubic.mm md, ms1       # output double widen
mqsubc.mm md, ms1       # output quadruple widen

# int matrix element-wise reverse subtract, md[i,j] = ms1[i,j] - md[i,j]
mrsubic.mm md, ms1
mwrsubc.mm md, ms1       # output double widen
mqrsubic.mm md, ms1      # output quadruple widen

# int matrix element-wise multiply with scalar int, md[i,j] = ms1[i,j] * rs2
memulc.mx md, ms1, rs2
mwemulc.mx md, ms1, rs2  # output double widen
mqemulc.mx md, ms1, rs2  # output quadruple widen

# int matrix element-wise multiply with int imm, md[i,j] = ms1[i,j] * imm
memulc.mi md, ms1, imm
mwemulc.mi md, ms1, imm  # output double widen
mqemulc.mi md, ms1, imm  # output quadruple widen

# float matrix element-wise add, md[i,j] = md[i,j] + ms1[i,j]
mfaddc.mm md, ms1
mfwadddc.mm md, ms1      # output double widen

# float matrix element-wise subtract, md[i,j] = md[i,j] - ms1[i,j]
mfsubc.mm md, ms1
mfwsubic.mm md, ms1      # output double widen
```

```

# float matrix element-wise reverse subtract, md[i,j] = ms1[i,j] - md[i,j]
mfrsubc.mm md, ms1
mfwrsubc.mm md, ms1      # output double widen

# float matrix element-wise multiply with scalar float, md[i,j] = ms1[i,j] * rs2
mfemulc.mf md, ms1, rs2
mfwemulc.mf md, ms1, rs2  # output double widen

```

4.4.3. Type-Convert Instructions

```

# convert float to float
mfncvtc.f.fw.m md, ms1      # double-width float to single-width float
mfwcvtc.fw.f.m md, ms1      # single-width float to double-width float

# convert int to float
mfecvtc.f.x.m md, ms1        # int to float

mfncvtc.f.xw.m md, ms1        # double-width int to float
mfncvtc.f.xq.m md, ms1        # quad-width int to float

mfwcvtc.fw.x.m md, ms1        # single-width int to double-width float
mfecvtc.fw.xw.m md, ms1        # double-width int to double-width float
mfncvtc.fw.xq.m md, ms1        # quad-width int to double-width float

# convert float to int
mfecvtc.x.f.m md, ms1        # float to int

mfwcvtc.xw.f.m md, ms1        # float to double-width int
mfwcvtc.xq.f.m md, ms1        # float to quad-width int

mfncvtc.x.fw.m md, ms1        # double-width float to single-width int
mfecvtc.xw.fw.m md, ms1        # double-width float to double-width int
mfwcvtc.xq.fw.m md, ms1        # double-width float to quad-width int

```

4.5. Instruction Listing

No.		31 28	27 20	19 15	14 12	11 7	6 0
Configuration		funct4		rs1	funct3	rd	opcode
1	msetypei	0000	mtypei[27:15]		111	rd	1110111
2	msetype	0001	00000000	rs1	111	rd	1110111
3	msettilemi	0010	mleni[27:15]		111	rd	1110111
4	msettilem	0011	00000000	rs1	111	rd	1110111

5	msettileki	0100	mleni[27:15]		111	rd	1110111	
6	msettilek	0101	00000000	rs1	111	rd	1110111	
7	msettileni	0110	mleni[27:15]		111	rd	1110111	
8	msettilen	0111	00000000	rs1	111	rd	1110111	
No.		31 26	25	24 20	19 15	14 12	11 7	6 0
Load		funct6	ls	rs2	rs1	eeew	md	opcode
9	mlae8.m	001001	0	rs2	rs1	000	md	1110111
10	mlae16.m	001001	0	rs2	rs1	001	md	1110111
11	mlae32.m	001001	0	rs2	rs1	010	md	1110111
12	mlae64.m	001001	0	rs2	rs1	011	md	1110111
13	mlbe8.m	001010	0	rs2	rs1	000	md	1110111
14	mlbe16.m	001010	0	rs2	rs1	001	md	1110111
15	mlbe32.m	001010	0	rs2	rs1	010	md	1110111
16	mlbe64.m	001010	0	rs2	rs1	011	md	1110111
17	mlce8.m	000000	0	rs2	rs1	000	md	1110111
18	mlce16.m	000000	0	rs2	rs1	001	md	1110111
19	mlce32.m	000000	0	rs2	rs1	010	md	1110111
20	mlce64.m	000000	0	rs2	rs1	011	md	1110111
9	mlate8.m	001101	0	rs2	rs1	000	md	1110111
10	mlate16.m	001101	0	rs2	rs1	001	md	1110111
11	mlate32.m	001101	0	rs2	rs1	010	md	1110111
12	mlate64.m	001101	0	rs2	rs1	011	md	1110111
13	mlbte8.m	001110	0	rs2	rs1	000	md	1110111
14	mlbte16.m	001110	0	rs2	rs1	001	md	1110111
15	mlbte32.m	001110	0	rs2	rs1	010	md	1110111
16	mlbte64.m	001110	0	rs2	rs1	011	md	1110111
17	mlcte8.m	000100	0	rs2	rs1	000	md	1110111
18	mlcte16.m	000100	0	rs2	rs1	001	md	1110111
19	mlcte32.m	000100	0	rs2	rs1	010	md	1110111

20	mlcte64.m	000100	0	rs2	rs1	011	md	1110111
Store		funct6	ls	rs2	rs1	ew	ms3	opcode
21	msae8.m	001001	1	rs2	rs1	000	ms3	1110111
22	msae16.m	001001	1	rs2	rs1	001	ms3	1110111
23	msae32.m	001001	1	rs2	rs1	010	ms3	1110111
24	msae64.m	001001	1	rs2	rs1	011	ms3	1110111
25	msbe8.m	001010	1	rs2	rs1	000	ms3	1110111
26	msbe16.m	001010	1	rs2	rs1	001	ms3	1110111
27	msbe32.m	001010	1	rs2	rs1	010	ms3	1110111
28	msbe64.m	001010	1	rs2	rs1	011	ms3	1110111
29	msce8.m	000000	1	rs2	rs1	000	ms3	1110111
30	msce16.m	000000	1	rs2	rs1	001	ms3	1110111
31	msce32.m	000000	1	rs2	rs1	010	ms3	1110111
32	msce64.m	000000	1	rs2	rs1	011	ms3	1110111
21	msate8.m	001101	1	rs2	rs1	000	ms3	1110111
22	msate16.m	001101	1	rs2	rs1	001	ms3	1110111
23	msate32.m	001101	1	rs2	rs1	010	ms3	1110111
24	msate64.m	001101	1	rs2	rs1	011	ms3	1110111
25	msbte8.m	001110	1	rs2	rs1	000	ms3	1110111
26	msbte16.m	001110	1	rs2	rs1	001	ms3	1110111
27	msbte32.m	001110	1	rs2	rs1	010	ms3	1110111
28	msbte64.m	001110	1	rs2	rs1	011	ms3	1110111
29	mscte8.m	000100	1	rs2	rs1	000	ms3	1110111
30	mscte16.m	000100	1	rs2	rs1	001	ms3	1110111
31	mscte32.m	000100	1	rs2	rs1	010	ms3	1110111
32	mscte64.m	000100	1	rs2	rs1	011	ms3	1110111
Arithmetic		funct6	fp	ms2	ms1	funct3	md	opcode
33	mma.mm	000000	0	ms2	ms1	110	md	1110111
34	mfma.mm	000000	1	ms2	ms1	110	md	1110111

35	mwma.mm	000001	0	ms2	ms1	110	md	1110111
36	mfwma.mm	000001	1	ms2	ms1	110	md	1110111
37	mqma.mm	000010	0	ms2	ms1	110	md	1110111
38	maddc.mm	000100	0	00000	ms1	110	md	1110111
39	mfaddc.mm	000100	1	00000	ms1	110	md	1110111
40	mwaddc.mm	000101	0	00000	ms1	110	md	1110111
41	mfwaddc.mm	000101	1	00000	ms1	110	md	1110111
42	mqaddc.mm	000110	0	00000	ms1	110	md	1110111
43	msubc.mm	000111	0	00000	ms1	110	md	1110111
44	mfsubc.mm	000111	1	00000	ms1	110	md	1110111
45	mwsubc.mm	001000	0	00000	ms1	110	md	1110111
46	mfwsubc.mm	001000	1	00000	ms1	110	md	1110111
47	mqsubc.mm	001001	0	00000	ms1	110	md	1110111
48	mrsubc.mm	001010	0	00000	ms1	110	md	1110111
49	mfrsubc.mm	001010	1	00000	ms1	110	md	1110111
50	mwrsubc.mm	001011	0	00000	ms1	110	md	1110111
51	mfwrsubc.mm	001011	1	00000	ms1	110	md	1110111
52	mqrsubc.mm	001100	0	00000	ms1	110	md	1110111
53	memulc.mx	001101	0	00000	ms1	110	md	1110111
54	mfemulc.mf	001101	1	00000	ms1	110	md	1110111
55	mwemulc.mx	001110	0	00000	ms1	110	md	1110111
56	mfwemulc.mf	001110	1	00000	ms1	110	md	1110111
57	mqemulc.mx	001111	0	00000	ms1	110	md	1110111
58	memulc.mi	010000	0	imm	ms1	110	md	1110111
59	mwemulc.mi	010001	0	imm	ms1	110	md	1110111
60	mqemulc.mi	010010	0	imm	ms1	110	md	1110111
Convert		funct6	fdst	ms2	ms1	funct3	md	opcode
61	mnfcvtc.f.fw.m	011000	1	00000	ms1	110	md	1110111
62	mwfcvtc.fw.f.m	011000	0	00000	ms1	110	md	1110111

63	mfcvtc.f.x.m	011010	1	00000	msl	110	md	1110111
64	mfcvtc.x.f.m	011010	0	00000	msl	110	md	1110111
65	mfncvtc.f.xw.m	011011	1	00000	msl	110	md	1110111
66	mfwcvtc.xw.f.m	011011	0	00000	msl	110	md	1110111
67	mfncvtc.f.xq.m	011100	1	00000	msl	110	md	1110111
68	mfwcvtc.xq.f.m	011100	0	00000	msl	110	md	1110111
69	mfwcvtc.fw.x.m	011101	1	00000	msl	110	md	1110111
70	mfncvtc.x.fw.m	011101	0	00000	msl	110	md	1110111
71	mfcvtc.fw.xw.m	011110	1	00000	msl	110	md	1110111
72	mfcvtc.xw.fw.m	011110	0	00000	msl	110	md	1110111
73	mfncvtc.fw.xq.m	011111	1	00000	msl	110	md	1110111
74	mfwcvtc.xq.fw.m	011111	0	00000	msl	110	md	1110111

Chapter 5. Intrinsic Examples

5.1. Matrix multiplication

```

void matmul_float16(c, a, b, m, k, n) {
    msettype(e16);                                // use 16bit input matrix element
    for (i=0; i<m; i+=tile_m) {                    // loop at dim m with tiling
        tile_m = msettile_m(m-i);
        for (j=0; j<n; j+=tile_n) {                // loop at dim n with tiling
            tile_n = msettile_n(n-j);

            acc = mfemul_mf(acc, 0.f)               // clear acc reg
            for (s=0; s<k; s+=tile_k) {             // loop at dim k with tiling
                tile_k = msettile_k(k-s);

                tr1 = mlae16_m(&a[i][s], k*2); // load left matrix a
                tr2 = mlbe16_m(&a[s][j], n*2); // load right matrix b
                acc = mfwma_mm(tr1, tr2);         // tiled matrix multiply,
                                                    // double widen output acc
            }

            acc = mfnvvt_f_fw_m(acc);              // convert widen result
            msce16_m(acc, &c[i][j], n*2);          // store to matrix c
        }
    }
}

```

5.2. Matrix multiplication with left matrix transposed

```

void matmul_a_tr_float16(c, a, b, m, k, n) {
    msettype(e16);                                // use 16bit input matrix element
    for (i=0; i<m; i+=tile_m) {                    // loop at dim m with tiling
        tile_m = msettile_m(m-i);
        for (j=0; j<n; j+=tile_n) {                // loop at dim n with tiling
            tile_n = msettile_n(n-j);

            acc = mfemul_mf(acc, 0.f)               // clear acc reg
            for (s=0; s<k; s+=tile_k) {             // loop at dim k with tiling
                tile_k = msettile_k(k-s);

                tr1 = mlate16_m(&a[s][i], m*2); // load transposed left matrix a
                tr2 = mlbe16_m(&a[s][j], n*2); // load right matrix b
                acc = mfwma_mm(tr1, tr2);         // tiled matrix multiply,
                                                    // double widen output acc
            }
        }
    }
}

```



```

        acc = mfncvt_f_fw_m(acc);           // convert widen result
        msce16_m(acc, &c[i][j], n*2);       // store to matrix c
    }
}

```

5.3. Matrix transpose without multiplication

```

void mattrans_float16(out, in, h, w) {
    msettype(e16);                               // use 16bit input matrix element

    for (i=0; i<h; i+=tile_m) {                   // loop at dim m with tiling
        tile_m = msettile_m(h-i);
        for (j=0; j<w; j+=tile_k) {               // loop at dim k with tiling
            tile_k = msettile_k(w-j);

            tr1 = mlae16_m(&in[i][j], w*2);       // load input matrix
            msate16_m(tr1, &out[j][i], h*2);      // store output matrix
        }
    }
}

```

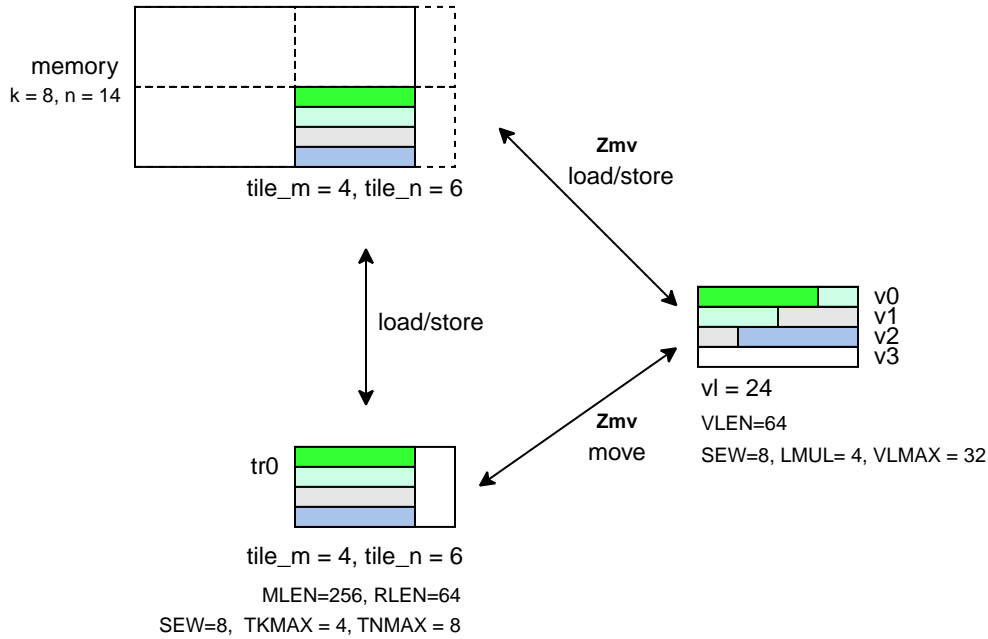
Chapter 6. Standard Matrix Extensions

6.1. Zmv: Matrix for Vector operations

The Zmv extension is defined to provide matrix support with the RISC-V Vector "V" extension.

The Zmv extension allows to load matrix tile slices into vector registers, and move data between slices of a matrix register and vector registers. Element-wise multiply between a matrix register and a vector register(broadcast to a matrix) is also supported.

The data layout examples of registers and memory in Zmv are shown below.



6.1.1. Load Instructions

```
# vd destination, rs1 base address, rs2 row byte stride
# lmul / (eew/sew) rows or columns

# for left matrix, a
mlae8.v    vd, (rs1), rs2 # 8-bit tile/acc slices load to vregs
mlae16.v   vd, (rs1), rs2 # 16-bit tile/acc slices load to vregs
mlae32.v   vd, (rs1), rs2 # 32-bit tile/acc slices load to vregs
mlae64.v   vd, (rs1), rs2 # 64-bit tile/acc slices load to vregs

# for right matrix, b
mlbe8.v    vd, (rs1), rs2 # 8-bit tile/acc slices load to vregs
mlbe16.v   vd, (rs1), rs2 # 16-bit tile/acc slices load to vregs
mlbe32.v   vd, (rs1), rs2 # 32-bit tile/acc slices load to vregs
```

```

mlbe64.v    vd, (rs1), rs2 # 64-bit tile/acc slices load to vregs

# for output matrix, c
mlce8.v     vd, (rs1), rs2 # 8-bit tile/acc slices load to vregs
mlce16.v    vd, (rs1), rs2 # 16-bit tile/acc slices load to vregs
mlce32.v    vd, (rs1), rs2 # 32-bit tile/acc slices load to vregs
mlce64.v    vd, (rs1), rs2 # 64-bit tile/acc slices load to vregs

```

6.1.2. Store Instructions

```

# vs3 store data, rs1 base address, rs2 row byte stride
# lmul / (eew/sew) rows or columns

# for left matrix, a
msae8.v     vs3, (rs1), rs2 # 8-bit tile/acc slices store from vregs
msae16.v    vs3, (rs1), rs2 # 16-bit tile/acc slices store from vregs
msae32.v    vs3, (rs1), rs2 # 32-bit tile/acc slices store from vregs
msae64.v    vs3, (rs1), rs2 # 64-bit tile/acc slices store from vregs

# for right matrix, b
msbe8.v     vs3, (rs1), rs2 # 8-bit tile/acc slices store from vregs
msbe16.v    vs3, (rs1), rs2 # 16-bit tile/acc slices store from vregs
msbe32.v    vs3, (rs1), rs2 # 32-bit tile/acc slices store from vregs
msbe64.v    vs3, (rs1), rs2 # 64-bit tile/acc slices store from vregs

# for output matrix, c
msce8.v     vs3, (rs1), rs2 # 8-bit tile/acc slices store from vregs
msce16.v    vs3, (rs1), rs2 # 16-bit tile/acc slices store from vregs
msce32.v    vs3, (rs1), rs2 # 32-bit tile/acc slices store from vregs
msce64.v    vs3, (rs1), rs2 # 64-bit tile/acc slices store from vregs

```

6.1.3. Data Move Instructions

Normal data move, rows or columns = lmul, eew = sew.

```

# Data move between matrix register rows and vector registers.

# vd[(i - rs2) * tile_n + j] = md[i, j], i = rs2 .. rs2 + lmul
mmvar.v.m   vd, ms1, rs2 # vl = lmul * tile_k
mmvbr.v.m   vd, ms1, rs2 # vl = lmul * tile_n
mmvcr.v.m   vd, ms1, rs2 # vl = lmul * tile_n

# md[i, j] = vd[(i - rs2) * tile_n + j], i = rs2 .. rs2 + lmul
mmvar.m.v   md, vs1, rs2 # vl = lmul * tile_k
mmvbr.m.v   md, vs1, rs2 # vl = lmul * tile_n
mmvcr.m.v   md, vs1, rs2 # vl = lmul * tile_n

```

```
# Data move between matrix register columns and vector registers.
```

```
# vd[(j - rs2) * tile_m + i] = md[i, j], j = rs2 .. rs2 + lmul
mmvac.v.m  vd, ms1, rs2  # vl = lmul * tile_m
mmvbc.v.m  vd, ms1, rs2  # vl = lmul * tile_k
mmvcc.v.m  vd, ms1, rs2  # vl = lmul * tile_m
```

```
# md[i, j] = vd[(j - rs2) * tile_m + i], j = rs2 .. rs2 + lmul
mmvac.m.v  md, vs1, rs2  # vl = lmul * tile_m
mmvbc.m.v  md, vs1, rs2  # vl = lmul * tile_k
mmvcc.m.v  md, vs1, rs2  # vl = lmul * tile_m
```

Double widen data move, rows or columns = $lmul / 2$, $eww = sew * 2$.

Only support accumulator registers as source and destination registers.

```
# Data move between matrix register rows and vector registers.
```

```
# vd[(i - rs2) * tile_n + j] = md[i, j], i = rs2 .. rs2 + lmul/2
mwmvcr.v.m vd, ms1, rs2  # vl = lmul * tile_n
```

```
# md[i, j] = vd[(i - rs2) * tile_n + j], i = rs2 .. rs2 + lmul/2
mwmvcr.m.v md, vs1, rs2  # vl = lmul * tile_n
```

```
# Data move between matrix register columns and vector registers.
```

```
# vd[(j - rs2) * tile_m + i] = md[i, j], j = rs2 .. rs2 + lmul/2
mwmvcc.v.m vd, ms1, rs2  # vl = lmul * tile_m
```

```
# md[i, j] = vd[(j - rs2) * tile_m + i], j = rs2 .. rs2 + lmul/2
mwmvcc.m.v md, vs1, rs2  # vl = lmul * tile_m
```

quadruple widen data move, rows or columns = $lmul / 4$, $eww = sew * 4$.

Only support accumulator registers as source and destination registers.

```
# Data move between matrix register rows and vector registers.
```

```
# vd[(i - rs2) * tile_n + j] = md[i, j], i = rs2 .. rs2 + lmul/4
mqmvcr.v.m vd, ms1, rs2  # vl = lmul * tile_n
```

```
# md[i, j] = vd[(i - rs2) * tile_n + j], i = rs2 .. rs2 + lmul/4
mqmvcr.m.v md, vs1, rs2  # vl = lmul * tile_n
```

```
# Data move between matrix register columns and vector registers.

# vd[(j - rs2) * tile_m + i] = md[i, j], j = rs2 .. rs2 + lmul/4
mqmvcc.v.m vd, ms1, rs2 # vl = lmul * tile_m

# md[i, j] = vd[(j - rs2) * tile_m + i], j = rs2 .. rs2 + lmul/4
mqmvcc.m.v md, vs1, rs2 # vl = lmul * tile_m
```

6.1.4. Matrix element-wise multiply

Matrix element-wise multiply instructions, the vector operand will broadcast to a matrix.

For m*emulcr.mv instructions, the vector operand will broadcast from a row to a c matrix.

```
# int matrix element-wise multiply with a row of vector int
# md[i,j] = ms1[i,j] * vs2[j]
memulcr.mv md, ms1, vs2
mwemulcr.mv md, ms1, vs2 # output double widen
mqemulcr.mv md, ms1, vs2 # output quadruple widen

# float matrix element-wise multiply with a row of vector float
# md = ms1[i,j] * vs2[j]
mfemulcr.mv md, ms1, vs2
mfwemulcr.mv md, ms1, vs2 # output double widen
```

For m*emulcc.mv instructions, the vector operand will broadcast from a column to a c matrix.

```
# int matrix element-wise multiply with a column of vector int,
# md[i,j] = ms1[i,j] * vs2[i]
memulcc.mv md, ms1, vs2
mwemulcc.mv md, ms1, vs2 # output double widen
mqemulcc.mv md, ms1, vs2 # output quadruple widen

# float matrix element-wise multiply with a column of vector float,
# md = ms1[i,j] * vs2[i]
mfemulcc.mv md, ms1, vs2
mfwemulcc.mv md, ms1, vs2 # output double widen
```

6.1.5. Instruction Listing

No.		31 26	25	24 20	19 15	14 12	11 7	6 0
	Load	funct6	ls	rs2	rs1	ew	md	opcode
1	mlae8.v	100000	0	rs2	rs1	000	md	1110111

No.		31 26	25	24 20	19 15	14 12	11 7	6 0
2	mlae16.v	100000	0	rs2	rs1	001	md	1110111
3	mlae32.v	100000	0	rs2	rs1	010	md	1110111
4	mlae64.v	100000	0	rs2	rs1	011	md	1110111
5	mlbe8.v	100001	0	rs2	rs1	000	md	1110111
6	mlbe16.v	100001	0	rs2	rs1	001	md	1110111
7	mlbe32.v	100001	0	rs2	rs1	010	md	1110111
8	mlbe64.v	100001	0	rs2	rs1	011	md	1110111
9	mlce8.v	100010	0	rs2	rs1	000	md	1110111
10	mlce16.v	100010	0	rs2	rs1	001	md	1110111
11	mlce32.v	100010	0	rs2	rs1	010	md	1110111
12	mlce64.v	100010	0	rs2	rs1	011	md	1110111
Store		funct6	ls	rs2	rs1	eev	ms3	opcode
13	msae8.v	100000	1	rs2	rs1	000	ms3	1110111
14	msae16.v	100000	1	rs2	rs1	001	ms3	1110111
15	msae32.v	100000	1	rs2	rs1	010	ms3	1110111
16	msae64.v	100000	1	rs2	rs1	011	ms3	1110111
17	msbe8.v	100001	1	rs2	rs1	000	ms3	1110111
18	msbe16.v	100001	1	rs2	rs1	001	ms3	1110111
19	msbe32.v	100001	1	rs2	rs1	010	ms3	1110111
20	msbe64.v	100001	1	rs2	rs1	011	ms3	1110111
21	msce8.v	100010	1	rs2	rs1	000	ms3	1110111
22	msce16.v	100010	1	rs2	rs1	001	ms3	1110111
23	msce32.v	100010	1	rs2	rs1	010	ms3	1110111
24	msce64.v	100010	1	rs2	rs1	011	ms3	1110111
Data Move		funct6	v2m	rs2	*s1	funct3	*d	opcode
25	mmvar.v.m	000001	0	rs2	ms1	101	vd	1110111
26	mmvar.m.v	000001	1	rs2	vs1	101	md	1110111
27	mmvbr.v.m	000010	0	rs2	ms1	101	vd	1110111

No.		31 26	25	24 20	19 15	14 12	11 7	6 0
28	mmvbr.m.v	000010	1	rs2	vs1	101	md	1110111
29	mmvcr.v.m	000011	0	rs2	ms1	101	vd	1110111
30	mmvcr.m.v	000011	1	rs2	vs1	101	md	1110111
31	mmvac.v.m	000101	0	rs2	ms1	101	vd	1110111
32	mmvac.m.v	000101	1	rs2	vs1	101	md	1110111
33	mmvbc.v.m	000110	0	rs2	ms1	101	vd	1110111
34	mmvbc.m.v	000110	1	rs2	vs1	101	md	1110111
35	mmvcc.v.m	000111	0	rs2	ms1	101	vd	1110111
36	mmvcc.m.v	000111	1	rs2	vs1	101	md	1110111
37	mwmvcr.v.m	010011	0	rs2	ms1	101	vd	1110111
38	mwmvcr.m.v	010011	1	rs2	vs1	101	md	1110111
39	mwmvcc.v.m	010111	0	rs2	ms1	101	vd	1110111
40	mwmvcc.m.v	010111	1	rs2	vs1	101	md	1110111
41	mqmvcr.v.m	100011	0	rs2	ms1	101	vd	1110111
42	mqmvcr.m.v	100011	1	rs2	vs1	101	md	1110111
43	mqmvcc.v.m	100111	0	rs2	ms1	101	vd	1110111
44	mqmvcc.m.v	100111	1	rs2	vs1	101	md	1110111
Arithmetic		funct6	fp	vs2	ms1	funct3	md	opcode
45	memulcr.mv	100001	0	vs2	ms1	110	md	1110111
46	mfemulcr.mv	100001	1	vs2	ms1	110	md	1110111
47	mwemulcr.mv	100010	0	vs2	ms1	110	md	1110111
48	mfwemulcr.mv	100010	1	vs2	ms1	110	md	1110111
49	mqemulcr.mv	100011	0	vs2	ms1	110	md	1110111
50	memulcc.mv	100100	0	vs2	ms1	110	md	1110111
51	mfemulcc.mv	100100	1	vs2	ms1	110	md	1110111
52	mwemulcc.mv	100101	0	vs2	ms1	110	md	1110111
53	mfwemulcc.mv	100101	1	vs2	ms1	110	md	1110111
54	mqemulcc.mv	100110	0	vs2	ms1	110	md	1110111

6.1.6. Intrinsic Examples: Matrix multiplication fused with element-wise vector operation

```

void fused_matmul_relu_float16(c, a, b, m, k, n) {
    msettype(e16);                                // use 16bit input matrix element
    for (i=0; i<m; i+=tile_m) {                    // loop at dim m with tiling
        tile_m = msettile_m(m-i);
        for (j=0; j<n; j+=tile_n) {                // loop at dim n with tiling
            tile_n = msettile_n(n-j);

            acc = mfemul_mf(acc, 0.f)               // clear acc reg
            for (s=0; s<k; s+=tile_k) {             // loop at dim k with tiling
                tile_k = msettile_k(k-s);

                tr1_r = mlae16_m(&a[i][s]);          // load left matrix a
                tr2_r = mlbe16_m(&a[s][j]);          // load right matrix b
                acc = mfwma_mm(tr1_r, tr2_r);        // tiled matrix multiply,
                                                    // double widen output acc
            }

            acc = mfnvvt_f_fw_m(acc);               // convert widen result to single

            for (s=0; s<tile_m; s+=8) {
                rows = min(tile_m - s, 8)
                vl = vsetvl(tile_n*rows, e16, m8)    // set vl to tile_n * rows

                v1 = mmvcr_v_m(acc_r[s])             // move acc rows to vreg
                v1 = vfmax_vf(0.f, v1)               // vfmax.vf for relu

                msce16_v(v1, &c[i+s][j], n);         // store output tile slices
            }
        }
    }
}

```

6.2. Zmbf16: Matrix Bfloat16(BF16) Extension

The Zmbf16 extension allows to use BF16 format as the data type of input matrix elements.

The Zmbf16 extension add a bit `mtype[4]` in `mtype` register.

Table 4. `mtype` register layout

Bits	Name	Description
XLEN-1	mill	Illegal value if set

Bits	Name	Description
XLEN-2:5	0	Reserved if non-zero
4	mbf16	Use BF16 input format
3	maccq	Support quad-width accumulator element
2:0	msew[2:0]	Selected element width (SEW) setting

The new **mtype** value is encoded in the immediate fields of msetypei, and in the rs1 register for msetype.

Suggested bf16 assembler name used for msetypei mtypei immediate

bf16 # Use BF16 format

Examples:

msetypei t0, e16, bf16 # SEW = 16, use BF16 as input matrix element

For implementation not support Bfloat16 format, **mtype.mill** will be set.

bf16 should be always used with **e16**(SEW=16), otherwise **mtype.mill** will be set.

6.3. Zmtf32: Matrix TensorFloat-32(TF32) Extension

The Zmtf32 extension allows to use TF32 FMA for matrix multiplication.

TF32 implementations are designed to achieve better performance on matrix multiplications and convolutions by rounding input Float32 data to have 10 bits of mantissa, and accumulating results with FP32 precision, maintaining FP32 dynamic range.

So when Zmtf32 is used, Float32 is still used as the input and output data type for matrix multiplication.

The Zmtf32 extension add a bit **mtype[5]** in **mtype** register.

Table 5. **mtype** register layout

Bits	Name	Description
XLEN-1	mill	Illegal value if set
XLEN-2:6	0	Reserved if non-zero
5	mtf32	Enable TF32 FMA for matrix multiplication
4	mbf16	Use bfloat16 input format

Bits	Name	Description
3	maccq	Support quad-width accumulator element
2:0	msew[2:0]	Selected element width (SEW) setting

The new **mtype** value is encoded in the immediate fields of msetypei, and in the rs1 register for msetype.

Suggested tf32 assembler name used for msetypei mtypei immediate

```
tf32 # enable TF32 FMA
```

Examples:

```
msetypei t0, e32, tf32 # SEW = 32, enable TF32 FMA
```

For implementation not support TF32 format, **mtype.mill** will be set.

tf32 should be always used with **e32**(SEW=32), otherwise **mtype.mill** will be set.

6.4. Zmic: Im2col Matrix Multiplication Extension

Im2col stands for Image to Column, and is an implementation technique of computing Convolution operation (in Machine Learning) using GEMM operations.

The Zmic extension allows to perform the im2col operation on-the-fly, by the new load instructions.

Working in progress.